Neuro-Symbolic AI: Bridging Learning and Reasoning for Next-Generation Intelligent Systems

Dhwani Singh
B.Tech Student
Department of Computer Science and Engineering
Delhi Technological University
(Formerly Delhi College of Engineering)
New Delhi, India
dhwanisingh_co21a3_66@dtu.ac.in

Druti Choudhury
B.Tech Student
Department of Computer Science and Engineering
Delhi Technological University
(Formerly Delhi College of Engineering)
New Delhi, India
drutichoudhury_c021a3_71@dtu.ac.in

Prof. Vinod Kumar
Professor
Department of Computer Science and Engineering
Delhi Technological University
(Formerly Delhi College of Engineering)
New Delhi, India
vinod_k@dtu.ac.in

***Abstract*—Neuro-Symbolic AI (NSAI) is an emerging field that seeks to combine the structured, rule-based reasoning of symbolic systems with the adaptive, data-driven learning capabilities of neural networks. This hybrid approach addresses the limitations of purely symbolic or purely neural models, creating AI systems that can reason about the world with the precision of logic while adapting to new data like neural models. By integrating these complementary paradigms, NSAI aims to develop more interpretable, scalable, and robust AI systems capable of tackling complex real-world problems.**

This literature review provides a comprehensive overview of the key components of NSAI, including foundational concepts like representation spaces, knowledge representation, and logical reasoning. It examines a range of neuro-symbolic architectures, from Logic Tensor Networks and Neuro-Symbolic Concept Learners to more advanced cognitive systems that incorporate both neural learning and symbolic logic. The paper also explores practical applications across diverse fields such as healthcare, geoscience, military systems, and patient monitoring, highlighting the transformative potential of NSAI. Additionally, it discusses the major challenges facing the field, including scalability, interpretability, knowledge acquisition, and ethical considerations, emphasizing the ongoing research needed to overcome these obstacles.

Furthermore, the review covers essential datasets and benchmarks like CLEVR, GQA, and NS-VQA, which play a critical role in evaluating the performance of neuro-symbolic systems. It also outlines promising future directions, including the development of unified representations, scalable learning algorithms, enhanced explainability, and ethically sound designs, which are essential for achieving human-level general intelligence.

Overall, this paper aims to provide researchers, practitioners, and industry leaders with a comprehensive understanding of the current landscape and future prospects of Neuro-Symbolic AI, offering a roadmap for the continued evolution of this critical field.

***Index Terms*—Neuro-Symbolic AI, Hybrid AI, Symbolic Reasoning, Neural Networks, Machine Learning, Artificial Intelligence, Cognitive Computing, Explainability, Scalability, Ethical AI, Knowledge Representation, Datasets and Benchmarks, Unified Representations, Future Directions, Cognitive Systems, Symbolic Logic, Neural-Symbolic Integration, Learning and Reasoning, Interpretable AI, Cognitive Architectures**

I. Introduction

The journey of artificial intelligence (AI) spans several decades, marked by shifting goals and evolving approaches. In its early days, AI pioneers believed that human cognition could be replicated through explicit, rule-based systems, a philosophy that came to define the **symbolic AI** era of the 1950s and 1960s. This period, often referred to as the **Good Old-Fashioned AI (GOFAI)** era, saw the development of some of the most iconic AI programs in history. Systems like **SHRDLU** for natural language understanding, **ELIZA** for psychotherapy simulation, **DENDRAL** for chemical structure analysis, and **MYCIN** for medical diagnosis were groundbreaking in their ability to mimic human reasoning through carefully crafted logic rules [1]. These systems could simulate conversations, draw chemical structures, and even diagnose medical conditions. However, their reliance on rigid, hand-coded rules made them brittle and inflexible, unable to cope with the complexities and uncertainties of real-world scenarios.

As the limits of purely symbolic approaches became apparent, researchers began exploring alternative methods. Real-world reasoning is rarely as straightforward as a set of fixed, predefined rules; it requires the ability to generalize, adapt, and learn from experience. This realization led to the emergence of **sub-symbolic** methods, which prioritize pattern recognition over explicit, logic-based reasoning. Early models like **Perceptrons** and **Hopfield Networks** laid the foundation for modern neural networks, which today drive major advances in areas like computer vision, natural language processing, and robotics. Contemporary architectures such as **Convolutional Neural Networks (CNNs)**, **Recurrent Neural Networks (RNNs)**, and **Transformers** (e.g., **BERT**, **GPT-4**) have achieved superhuman performance in tasks ranging from image recognition to language translation [1]. However, despite their remarkable capabilities, these systems often act as "black boxes," making their internal decision-making processes difficult to interpret.

A. Bridging the Gap: The Need for Neuro-Symbolic AI

Despite the breakthroughs in sub-symbolic AI, purely neural systems face several fundamental challenges:

- **Knowledge Representation and Reasoning:** While neural networks excel at identifying statistical patterns in raw data, they struggle to capture complex, hierarchical relationships. For instance, a medical AI system must understand that "heart disease" is a type of "cardiovascular condition," which in turn is a subclass of "diseases" [2]. Symbolic systems are inherently better at handling such structured, taxonomic relationships.

- **Compositionality and Generalization:** Human reasoning is inherently compositional, allowing us to combine familiar concepts in novel ways. For example, understanding the phrase "a red, round object on a blue table" requires integrating multiple attributes simultaneously. Neural models often require vast amounts of data to achieve this level of flexible reasoning, while symbolic systems can handle it more naturally [2].

- **Explainability and Trustworthiness:** In critical fields like healthcare, finance, and autonomous systems, the ability to explain and justify decisions is essential. Symbolic methods provide clear, traceable decision paths, while neural networks, despite their power, often lack this level of interpretability [2].

B. The Rise of Neuro-Symbolic AI

To address these limitations, researchers have developed **Neuro-Symbolic AI**, a hybrid approach that combines the best of both worlds:

- **Structured Knowledge Integration:** Neuro-Symbolic systems leverage the structured, rule-based reasoning of symbolic methods alongside the pattern recognition strengths of neural networks, enabling more sophisticated decision-making without sacrificing flexibility.

- **Efficient Learning with Prior Knowledge:** Unlike purely neural models that require extensive training data, Neuro-Symbolic approaches can incorporate existing domain knowledge, reducing the amount of data needed for effective learning.

- **Interpretability and Transparency:** By integrating symbolic reasoning with statistical learning, Neuro-Symbolic systems offer clearer, more interpretable decision processes, making them particularly suitable for regulated industries where transparency is crucial.

C. Structure and Scope of the Paper

This paper is structured as follows. **Section II** traces the historical evolution of Neuro-Symbolic AI, from its early symbolic roots to modern hybrid approaches. **Section III** covers the key concepts that form the theoretical backbone of this field, including knowledge representation and reasoning. **Section IV** discusses the major architectures and approaches that define Neuro-Symbolic AI. **Section V** explores real-world applications, while **Section VI** examines the challenges and open problems that remain. **Section VII** reviews commonly used datasets and benchmarks, and **Section VIII** highlights future directions. Finally, **Section IX** provides a summary.

II. Historical Background

The history of Neuro-Symbolic AI traces back to the foundational ideas of artificial intelligence itself, blending the structured, rule-based logic of early symbolic systems with the adaptive, data-driven learning of modern neural networks. This journey reflects the shifting priorities of AI research, moving from purely logical approaches to more flexible, pattern-based methods as the complexity of real-world problems became evident.

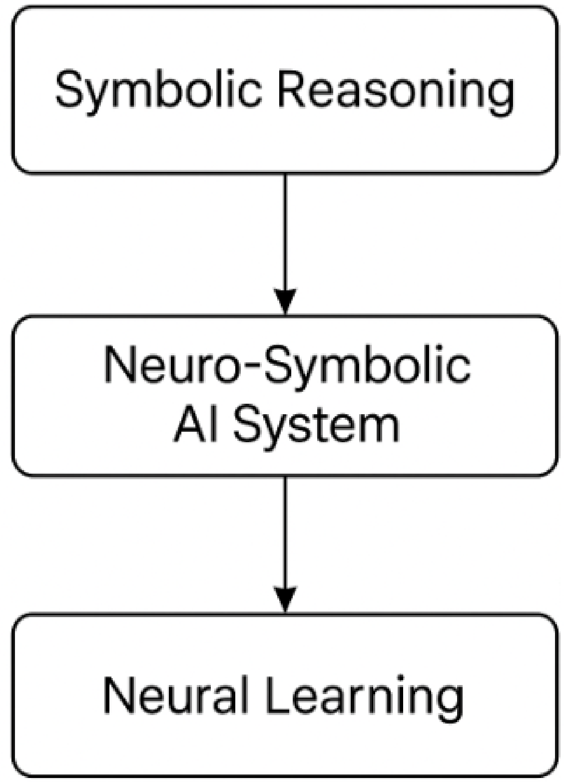

Fig. 1: High-Level Architecture of a Neuro-Symbolic AI System, integrating symbolic reasoning with neural learning (self-created).

A. Early Symbolic AI (1950s - 1980s)

Symbolic AI, often referred to as **Good Old-Fashioned AI (GOFAI)**, emerged in the 1950s, characterized by systems that relied on formal logic, rule-based reasoning, and explicit knowledge representation. These early efforts aimed to replicate human thinking through structured, human-readable symbols and predefined rules, making them inherently interpretable. Pioneering examples include **SHRDLU** (Winograd, 1971), a system capable of understanding and manipulating blocks through natural language, **ELIZA** (Weizenbaum, 1966), a program simulating psychotherapy conversations, and **DENDRAL** (Lindsay et al., 1980), a rule-based expert system for chemical analysis [1].

However, despite their impressive early successes, these systems struggled with several critical challenges. They required extensive manual encoding of domain-specific rules, making them brittle and difficult to scale. Moreover, symbolic systems were hampered by the **Frame Problem** (McCarthy, 1969) – the difficulty of specifying all relevant knowledge needed to respond appropriately in dynamic environments – and the broader **Commonsense Knowledge Problem**, which highlighted their inability to handle the vast, implicit background knowledge humans use effortlessly [3].

B. Rise of Sub-Symbolic AI (1980s - 2000s)

In contrast, the sub-symbolic approach emerged in the 1980s, inspired by biological models of neural processing. This wave, driven by the resurgence of **Artificial Neural Networks (ANNs)**, emphasized learning directly from data without the need for explicit symbolic representations. Foundational breakthroughs included the **Perceptron** (Rosenblatt, 1958), which demonstrated that simple neuron-like structures could learn linearly separable patterns, the introduction of **Backpropagation** (Rumelhart et al., 1986), which enabled multi-layer networks to learn complex functions, and **Hopfield Networks** (Hopfield, 1982), which modeled associative memory [4].

While these approaches excelled at pattern recognition and statistical learning, they introduced a new set of challenges. Chief among these was the **Black Box Problem** – the difficulty in interpreting how these networks make decisions. Unlike symbolic systems, which are inherently explainable, neural networks often lack transparency, making it hard to understand their reasoning, a critical limitation in fields like healthcare and autonomous systems where accountability is essential [5].

C. Early Hybrid Systems (1990s - 2010s)

Recognizing the limitations of purely symbolic and purely sub-symbolic approaches, researchers began exploring hybrid systems as early as the 1990s. These efforts sought to combine the structured, rule-based reasoning of symbolic AI with the adaptive learning capabilities of neural networks. Notable early work included **Hybrid Logic Networks** and **Logic Tensor Networks (LTNs)**, which integrated symbolic logic directly into neural architectures, allowing for richer, more interpretable learning models [2].

During this period, significant progress was also made in combining symbolic reasoning with probabilistic models, resulting in frameworks like **Probabilistic Logic Networks** and **Markov Logic Networks (MLNs)**. These systems bridged the gap between structured logic and statistical reasoning, providing a foundation for modern neuro-symbolic systems [4].

D. Modern Neuro-Symbolic AI (2010s - Present)

The past decade has seen a resurgence of interest in neuro-symbolic approaches, driven by the demand for AI systems that can combine the flexibility of neural learning with the transparency and structure of symbolic reasoning. Key innovations include **DeepProbLog** (Manhaeve et al., 2018), which integrates probabilistic logic programming with deep learning, **Neural Theorem Provers (NTPs)** (Rocktäschel and Riedel, 2017), which enable neural networks to perform symbolic reasoning, and **Neural Logic Machines (NLMs)** (Dong et al., 2019), which extend these ideas to relational reasoning [3].

More recent advances have focused on creating systems capable of reasoning over complex, unstructured data, bridging the gap between symbolic semantics and sub-symbolic learning. This includes approaches like **Semantic Loss Networks**, which use logic as a regularization term to guide neural training, and **Graph Neural Networks (GNNs)**, which provide powerful tools for relational reasoning [4]. These systems represent the cutting edge of neuro-symbolic AI, pushing the field closer to truly general, human-like intelligence.

III. Key Concepts and Foundations

Neuro-Symbolic AI (NSAI) represents a forward-looking approach that blends the structured, rule-based logic of symbolic systems with the adaptive, data-driven learning capabilities of neural networks. This combination aims to bridge the gap between human-like understanding and the powerful pattern recognition capabilities of modern AI systems. Understanding the foundational elements of NSAI is crucial for appreciating its potential, including representation spaces, knowledge representation, reasoning mechanisms, the critical role of logic, and the unique challenges these systems face as they evolve.

A. Representation Spaces

Representation spaces serve as the underlying frameworks within NSAI, defining how information is encoded, transformed, and processed within an AI system. These spaces are critical because they allow the structured logic of symbolic reasoning to coexist with the flexible, pattern-based learning of neural networks. According to Zhang and Sheng (2024), these representation spaces can be broadly categorized based on their ability to integrate symbolic and neural data: [6]

- **Single-modal, non-heterogeneous**: Designed to handle a single type of data, like text or images, without integrating symbolic and neural elements. These systems typically excel at specialized tasks where the data is uniform and well-defined.
- **Multi-modal, non-heterogeneous**: Capable of processing multiple data types independently, without true integration between symbolic and neural components. This approach is often used in multimedia systems where different data streams are processed separately.
- **Single-modal, heterogeneous**: Focused on a single data type but integrates symbolic and neural representations within a unified framework, creating a more comprehensive understanding of the data.
- **Multi-modal, heterogeneous**: Processes diverse data types while integrating symbolic and neural representations, allowing for richer, context-aware data processing. These systems can handle complex tasks where multiple data types must be analyzed simultaneously.

- **Dynamic adaptive**: Adapts its representation strategies based on the task at hand, providing flexibility and scalability. This approach is particularly useful in environments with rapidly changing data and complex decision-making requirements.

These categories highlight the growing sophistication of modern NSAI systems, which must balance interpretability, scalability, and task-specific flexibility to achieve robust performance. [6]

	Single Modal Data	Multi-Modal Data	Neural Network OR Symbolic Logic Representation	Neural Network AND Symbolic Logic Representation
Single-modal and non-heterogeneous	✓		✓	
Multimodal and non-heterogeneous		✓	✓	
Single-modal and heterogeneous	✓			✓
Multimodal and heterogeneous		✓		✓
Dynamic adaptive	✓	✓	✓	✓

Fig. 2: Neuro-Symbolic AI's classification by representation space (Zhang and Sheng, 2024) [6].

B. Knowledge Representation

Knowledge representation (KR) is a cornerstone of Neuro-Symbolic AI (NSAI), providing the critical link between low-level neural activations and high-level symbolic reasoning. It enables the integration of discrete, human-readable symbolic structures with continuous, data-driven neural representations, forming the backbone of intelligent systems that can both learn from data and reason logically. [7]

1) Basic Structure and Challenges of Knowledge Representation: Knowledge representation in NSAI acts as a bridge, connecting the symbolic world of structured logic with the adaptive, flexible nature of neural networks. This connection allows for complex reasoning by combining the clarity and interpretability of symbolic structures with the data-driven learning of neural components. [7]

2) Two Primary Integration Approaches:

a) Representational Integrations: These approaches tightly couple symbolic logic and neural learning, creating unified knowledge bases that leverage the strengths of both paradigms. This style of integration aims to overcome the limitations of purely connectionist approaches by combining the precise, interpretable reasoning of symbolic systems with the flexible, adaptive learning of neural networks.

b) Modular Integrations: In contrast, modular integrations maintain a clear separation between symbolic and neural components, allowing for independent processing and well-defined interfaces. This design promotes flexibility, scalability, and transparency, making it ideal for complex systems where different components must work together without being tightly coupled. [7]

3) Functional Approaches within Modular Systems: Modular integrations can be further categorized based on the degree of interconnection and the flow of information:

- **Passively Coupled:** Systems with loose connections, typically communicating through shared files or external data transfers, allowing for independent operation.
- **Actively Coupled:** Tighter integration with shared memory and synchronized processing, facilitating more direct and responsive interactions.
- **Interleaved:** High-level interactions that involve direct function calls and complex communication protocols, enabling real-time information exchange. [7]

4) Subcategories of Representational Integrations: Within representational integrations, some systems place more emphasis on the symbolic component, using neural methods to enhance specific tasks, while others prioritize the neural component, incorporating symbolic rules as supporting structures within predominantly neural frameworks. [7]

5) Advantages of Knowledge Representation in NSAI: Symbolic knowledge representation provides clarity, modularity, and transparency, making it particularly well-suited for applications that require detailed explanations and structured reasoning. Neural components, on the other hand, offer flexibility, efficiency, and the ability to learn directly from data, complementing the structured reasoning of symbolic systems to create more robust, intelligent models. [7]

C. The Role of Logic in Neuro-Symbolic Systems

Logic plays a foundational role in NSAI, providing the structure needed for complex reasoning and decision-making. Belle (2024) argues that symbolic logic, despite its age, remains highly relevant for building robust AI systems. He emphasizes that symbolic logic can capture uncertain and probabilistic relationships, a critical capability for AI systems operating in real-world environments filled with ambiguity and incomplete information [8].

Logical frameworks also support explainability, a key requirement in high-stakes applications like healthcare, finance, and autonomous systems. Unlike purely neural models, which often function as ”black boxes,” logic-based systems offer clear, human-readable decision paths, making them essential for applications where accountability is critical [8].

D. Challenges and Future Directions

Despite significant progress, integrating symbolic and neural methods remains challenging. Key obstacles include:

- **Scalability**: As the complexity of symbolic reasoning grows, maintaining efficiency becomes increasingly difficult.
- **Knowledge Acquisition Bottleneck**: Encoding human knowledge in symbolic form is often labor-intensive and error-prone.
- **Interpretability vs. Flexibility Trade-off**: Balancing the interpretability of symbolic systems with the flexibility of neural networks remains a persistent challenge.

However, ongoing research in areas like differentiable logic, neural-symbolic architectures, and automated knowledge extraction promises to address these limitations, pushing NSAI closer to achieving human-level general intelligence.

IV. Approaches and Architectures

Neuro-Symbolic AI (NSAI) seeks to bridge the gap between symbolic reasoning, known for its clarity and precision, and the flexible, data-driven learning capabilities of neural networks. This hybrid approach aims to combine the best of both paradigms, creating AI systems that can reason about the world with the precision of logic while adapting to new data like neural models. This section explores a range of approaches and architectures within NSAI, each tailored to specific challenges in knowledge representation, learning, and reasoning.

A. Symbolic-Neural Integration

Symbolic-neural integration forms the backbone of many NSAI systems, tightly coupling logical reasoning with neural learning to create unified models. One notable example is the Logic Tensor Network (LTN) framework developed by Serafini and Garcez (2016). LTNs directly incorporate first-order logic into neural architectures, enabling them to process both structured, symbolic data and continuous inputs. This combination supports a wide range of applications, including natural language understanding, autonomous robotics, and scientific data analysis, where precise reasoning and flexible learning are critical. [9]

Another influential approach is the Neuro-Symbolic Concept Learner (NSCL), which integrates neural perception with explicit symbolic reasoning. These models excel in complex tasks like visual question answering (VQA), where the ability to reason about the relationships between objects in a scene is essential. For instance, NSCLs have been benchmarked on datasets like CLEVR and GQA, demonstrating superior performance in structured reasoning compared to purely neural models. [10] [11]

Recent advancements in this area have also included probabilistic models like DeepProbLog, which extend symbolic logic to handle uncertainty, making them suitable for real-world applications like financial forecasting, medical diagnosis, and robotics. These models blend probabilistic reasoning with deep learning, allowing systems to manage ambiguous or noisy inputs effectively. [12] [13]

B. Modular and Distributed Architectures

Modular architectures maintain a clear separation between symbolic and neural components, promoting transparency and scalability. In these systems, the symbolic reasoning engine and neural perception modules operate independently, communicating through well-defined interfaces. This design is particularly advantageous in systems requiring clear explanations and high interpretability. For example, Paul et al. (2024) developed a modular framework that decouples symbolic reasoning from perceptual data processing, improving both transparency and adaptability. [14]

Distributed architectures extend this concept further, enabling symbolic and neural components to operate across multiple computational nodes. This approach supports large-scale, real-time processing, making it ideal for applications like autonomous driving, smart city management, and large-scale financial systems. Foundation models like GPT-4 have been adapted to these architectures, integrating pre-trained knowledge bases for rapid, context-aware reasoning. [13] [12]

Additionally, the development of neurosymbolic processors and specialized AI chips has significantly improved the efficiency of these systems, supporting high-speed, large-scale reasoning across distributed networks. These advancements enable real-time decision-making in critical applications, such as industrial automation and defense systems. [12]

C. End-to-End Differentiable Architectures

End-to-end differentiable architectures directly integrate symbolic reasoning components within neural networks, allowing the entire system to be optimized through gradient descent. This design is exemplified by Logic Tensor Networks (LTNs) and Neuro-Symbolic Concept Learners (NSCLs), which jointly optimize symbolic and neural parameters, enabling precise, context-aware reasoning. [9] [12]

These architectures are particularly effective in robotics and autonomous systems, where continuous learning is essential for adapting to changing environments. For instance, autonomous drones and robotic surgery systems benefit from this approach, as they must make real-time decisions based on both learned experiences and explicit rules. [12]

D. Neural-Symbolic Memory Systems

Neural-symbolic memory systems store and retrieve structured symbolic information while leveraging the generalization capabilities of neural networks. These systems incorporate memory modules that explicitly represent symbolic knowledge, enabling long-term reasoning and complex multi-step decision-making. Benchmarks like CLEVR and GQA are often used to evaluate these systems, as they require both visual perception and structured reasoning. [10] [11]

Recent innovations have included episodic memory systems, which allow AI to retain and recall past experiences over extended periods, supporting more sophisticated decision-making in dynamic environments. This capability is critical for applications like military strategy planning, disaster response, and intelligent traffic management. [13] [15]

E. Cognitive Architectures for Neuro-Symbolic AI

Cognitive architectures aim to replicate human-like reasoning by integrating symbolic logic with neural learning in a biologically plausible manner. These systems often incorporate working memory, episodic memory, and attention mechanisms to simulate complex cognitive processes. Wan et al. (2024) propose a cognitive architecture that combines symbolic reasoning with neural pattern recognition, providing a more holistic approach to AI. [12]

Cunnington et al. (2024) also emphasize the role of foundation models in cognitive architectures, highlighting their ability to integrate vast amounts of symbolic knowledge with deep learning for enhanced context-awareness and long-term reasoning. This integration is critical for developing truly intelligent systems capable of understanding and responding to complex, real-world situations. [13]

F. Challenges in Neuro-Symbolic Architectures

Despite significant progress, several challenges remain in designing effective neuro-symbolic architectures:

- **Scalability:** As the complexity of symbolic reasoning increases, maintaining efficiency becomes challenging.
- **Interoperability:** Ensuring smooth communication between symbolic and neural components remains a critical hurdle.
- **Explainability vs. Flexibility:** Balancing interpretability with the flexibility of neural networks is a persistent issue.
- **Knowledge Acquisition Bottlenecks:** Encoding human knowledge into symbolic form is labor-intensive and error-prone.

Ongoing research in areas like differentiable logic, automated knowledge extraction, neurosymbolic hardware acceleration, cognitive modeling, and efficient workload characterization aims to address these challenges, pushing NSAI closer to achieving human-level general intelligence. [12] [15]

V. Applications of Neuro-Symbolic AI

Neuro-Symbolic AI (NSAI) combines the learning flexibility of neural networks with the interpretability and reasoning power of symbolic systems, showing great promise in a variety of real-world applications. This section offers a thorough analysis of its uses in patient monitoring, healthcare, geoscience, and military operations, emphasising the revolutionary potential of this hybrid method.

A. Healthcare Applications

In the healthcare industry, where interpretable, dependable, and effective AI systems are essential, NSAI is becoming more and more significant. A thorough analysis of NSAI's effects on healthcare is given by Hossain and Chen (2025), who emphasise how it can facilitate personalised therapy, increase diagnostic precision, and improve decision-making. One important field is drug discovery, where NSAI models outperform conventional machine learning techniques in identifying possible drug candidates by fusing biological knowledge with data-driven learning. Symbolic reasoning should be used to find patterns and the connections in intricate biomedical data, these systems offer insights that merely data-driven models sometimes overlook. Moreover, NSAI can speed up the development of novel therapeutic proteins and pharmacological compounds by fusing domain-specific knowledge with high-throughput experimental data. [16]

Additionally, NSAI has demonstrated efficacy in automating early disease identification and patient monitoring. Logic Tensor Networks (LTNs), for instance, have been used to forecast the course of diseases, find early warning indicators, and spot anomalies in patient data. This integration of symbolic logic with neural learning allows healthcare systems to make sense of complex, noisy data while maintaining interpretability, which is essential for clinical decision support systems. [16]

B. Geoscience and Mineral Prediction

In geoscience, NSAI has been utilized to address challenges in mineral exploration and resource management. Chen et al. (2024) highlight how NSAI systems can improve the accuracy of mineral prediction by integrating geological domain knowledge with data-driven learning. These systems can analyze large, heterogeneous datasets, including geochemical assays, remote sensing images, and geological maps, to identify potential mineral-rich zones. Unlike conventional machine learning methods, which often struggle with sparse and noisy data, NSAI can incorporate expert knowledge through symbolic representations, enhancing model interpretability and robustness. [16]

Additionally, NSAI systems can model complex geological processes, such as the formation of mineral deposits, by combining physical rules with data-driven insights. This approach not only improves prediction accuracy but also provides geologists with clearer explanations of the underlying factors driving mineralization, which is crucial for making informed exploration decisions. [16]

C. Military Applications

The revolutionary potential of NSAI in military settings, where danger detection, autonomous system control, and real-time decision-making are crucial, is covered by Hagos and Rawat (2024). By combining several data streams—such as sensor inputs, satellite imagery, and battlefield reports—into a logical operational picture, these systems can improve situational awareness. NSAI can be used to predict possible threats, automate the processing of satellite photos for enemy movement identification, and provide military leaders with real-time recommendations for the best course of action. These systems' symbolic element lowers the possibility of unforeseen repercussions in high-stakes situations by guaranteeing that judgements are auditable and explicable. [17]

Additionally, NSAI is used in cybersecurity, battlefield logistics, and autonomous weapon systems. Neural networks and symbolic reasoning, for example, can be combined to create stronger cybersecurity defences that can recognise and neutralise online threats instantly. This hybrid method also facilitates the creation of intelligent training systems, which can mimic complex combat scenarios to boost military preparedness and operational effectiveness. [17]

D. Patient Monitoring and Human Activity Recognition

Another crucial use of NSAI is patient monitoring, where real-time data analysis is necessary to guarantee patient safety and care quality. In order to identify essential health events, Fenske et al. (2024) present a system that integrates sensor data with symbolic domain knowledge to monitor patient behaviours in hospitals using NSAI. This method, for instance, can be used to measure patient movement in real time, track vital signs, and identify falls, giving medical personnel useful information to enhance patient outcomes. Unlike purely data-driven methods, NSAI systems can integrate expert knowledge and medical advice to consistently identify critical events even in noisy, sensor-rich environments. [18]

These systems also address key challenges in patient monitoring, such as data heterogeneity and the need for low-latency processing. By combining neural networks for perception with symbolic reasoning for high-level understanding, NSAI provides a more comprehensive and interpretable solution for patient care, reducing the risk of false alarms and enhancing clinical decision-making. [18]

E. Future Directions and Emerging Applications

As NSAI continues to evolve, its applications are expanding into other critical areas, including finance, autonomous transportation, natural language processing, and robotics. Ongoing research aims to refine these systems to improve scalability, interpretability, and efficiency, ensuring they can meet the demands of increasingly complex real-world tasks.

VI. Challenges in Neuro-Symbolic AI

By combining the adaptive, data-driven learning of neural networks with the structured, rule-based reasoning of symbolic systems, neuro-symbolic artificial intelligence (NSAI) provides a potent framework. However, in order to reach its full potential, this hybrid approach also presents a number of difficult problems that need to be resolved. These difficulties represent the essential distinctions between neural approaches and symbolic approaches to artificial intelligence and cut beyond the technical, operational, and ethical spheres. The main issues raised by current studies are covered in this part, including knowledge acquisition, scalability, interpretability, robustness, and ethical issues.

A. Scalability and Efficiency

For NSAI systems, which frequently need significant computational resources to successfully combine discrete symbolic reasoning with dense, continuous neural processing, scalability is still a major difficulty. The sparse, rule-based operations typical of symbolic thinking are better suited for traditional neural accelerators, such as GPUs and TPUs, which are optimised for dense matrix operations. This discrepancy can limit the real-time performance of NSAI systems in complex situations and lead to major inefficiencies. [15]

Researchers are looking on specialised hardware solutions that can effectively handle both the symbolic and neural workloads in order to overcome these constraints. This includes the creation of neurosymbolic processors, that are intended to close the gap between discrete reasoning and continuous learning and may allow for real-time decision-making in the applications involving large amounts of data. [15]

B. Explainability and Interpretability

The intrinsic transparency of symbolic reasoning is one of its main benefits; this is sometimes lost when paired with the more opaque, data-driven neural models. In crucial applications like healthcare, banking, and autonomous systems, where transparent, intelligible decisions are crucial for fostering confidence and guaranteeing regulatory compliance, this lack of interpretability poses serious difficulties. [19]

The development of hybrid models that retain both neural and explicit symbolic reasoning, as well as techniques for deriving human-readable rules from trained neural networks, are attempts to overcome this difficulty. It is still difficult to strike a balance between performance and interpretability, though, because increasing transparency may limit the model's flexibility and learning potential. [19] [12]

C. Robustness and Generalization

Although this is still a constant problem, NSAI systems need to be reliable and able to generalise to new, unknown data. Reliability is crucial in adversarial situations, when even slight changes in input data can have a big influence on system performance, according to Hagos and Rawat (2024). [17]

Furthermore, NSAI models may have trouble adapting to changes in the distribution of data, which makes it challenging to guarantee consistent behaviour in a variety of settings. Because autonomous systems frequently deal with erratic inputs and quickly shifting circumstances, this problem is very important. [17] [12]

D. Knowledge Acquisition and Representation

The development of NSAI systems is severely hampered by the difficulty and error-proneness of translating human expertise into symbolic form. This procedure frequently necessitates in-depth topic expertise and exact validation because even small errors might result in significant system faults. [20]

Furthermore, there are constant difficulties in preserving coherence between structured symbolic information and learnt neural representations. The stability of the entire system can be jeopardized by symbolic rules, which are usually strict and need to be updated frequently when new information becomes available. [12]

E. Ethical and Strategic Considerations

Because automated decision-making can have serious human repercussions in high-stakes applications like healthcare and military systems, ethical problems are especially acute. Hagos and Rawat (2024) draw attention to the strategic dangers of relying too much on autonomous systems, such as the probability of abuse, the lack of human supervision, and the increasing hostilities. [17]

Strong security measures are becoming increasingly necessary as NSAI systems become more powerful to stop misuse and illegal access. This entails protecting private information and making sure self-governing systems follow accepted moral standards. [17] [20]

F. Ongoing Research and Future Directions

A variety of approaches are being actively investigated by researchers to address these issues, including as the creation of more effective hardware architectures, enhanced interpretability strategies, and automated knowledge extraction processes. Furthermore, work is being done to develop stronger training methods that can manage hostile and noisy data more effectively, improving the dependability of NSAI systems in practical applications. [15] [19] [12]

Overcoming these obstacles will be necessary to reach general intelligence on par with that of humans and to ensure that these systems can function securely and efficiently in a variety of real-world settings as NSAI develops.

VII. Datasets and Benchmarks for Neuro-Symbolic AI

Neuro-Symbolic AI (NSAI) systems require specialized datasets and benchmarks to evaluate their reasoning, perception, and language understanding capabilities. Unlike traditional machine learning datasets, these benchmarks are designed to test both symbolic reasoning and neural perception, ensuring that models can effectively integrate structured logic with data-driven learning. Three significant datasets have emerged as critical benchmarks in this domain: **CLEVR**, **GQA**, and **NS-VQA**, each offering unique challenges and insights into the performance of NSAI systems.

A. CLEVR Dataset

The CLEVR (Compositional Language and Elementary Visual Reasoning) as a synthetic benchmark was designed to assess the visual reasoning abilities of AI models. It features 100,000 photos and more than a million computer generated questions, along with a variety of reasoning exercises such attribute recognition, spatial correlations, counting, logical comparisons, and multi-step reasoning. [10]

The main features of the CLEVR dataset are:

- **Synthetic Scenes**: The dataset uses simple, 3D-rendered objects with well-defined attributes, including color, shape, size, and material. This controlled environment eliminates the ambiguity often present in natural images, allowing precise evaluation of reasoning capabilities.
- **Functional Programs**: Each question in CLEVR is represented as a functional program, detailing the reasoning steps required to arrive at the correct answer. This structure enables models to execute complex reasoning sequences and provides a clear evaluation framework.
- **Bias Reduction**: CLEVR minimizes question-conditional biases through rejection sampling and structured data generation, ensuring that models cannot rely on statistical shortcuts to solve tasks without genuine reasoning.
- **Hierarchical Reasoning**: The dataset is specifically designed to test hierarchical reasoning, requiring models to understand multi-step dependencies and complex relational logic. [10]

Basic attribute identification questions (like *What color is the cylinder?*) and more intricate compositional queries (like *Are there more red cubes than blue spheres?*) are examples of the types of questions that are part of CLEVR. [10]

B. GQA Dataset

The GQA (Generalized Question Answering) dataset, developed by Hudson and Manning (2019), extends the ideas of CLEVR to real-world scenes, introducing more diverse and challenging reasoning tasks. It includes over 22 million questions grounded in natural images, significantly increasing the complexity of visual understanding tasks. [11]

Key characteristics of the GQA dataset include:

- **Real-World Images**: Unlike CLEVR, which uses synthetic data, GQA is built on real-world photographs from the Visual Genome dataset, providing a richer, more realistic testing environment.
- **Scene Graph Annotations**: The logical structure and semantics of the queries can be carefully controlled since each image is represented as a structured scene graph that documents objects, properties, and relationships.
- **Compositional Reasoning**: GQA questions are generated using structured templates that promote compositional reasoning, requiring models to interpret complex spatial and semantic relationships.
- **Bias Mitigation**: The dataset employs sophisticated techniques to reduce statistical biases, including tunable smoothing methods that balance answer distributions and minimize overfitting to common patterns. [11]
- **Extensive Evaluation Metrics**: GQA introduces new performance metrics, including consistency, grounding, and plausibility, providing a more nuanced assessment of model capabilities beyond simple accuracy. [11]

The GQA dataset tests models' capacity to combine linguistic patterns, object relationships, and scene context into a cohesive framework, pushing them beyond basic object recognition.

C. NS-VQA Dataset

The NS-VQA (Neural-Symbolic Visual Question Answering) dataset, presented by Yi et al. (2019), builds on the CLEVR framework but introduces a more explicitly structured approach to reasoning. It is specifically designed to disentangle vision and reasoning, providing clear semantic representations that separate perception from logic. [21]

Key features of the NS-VQA dataset include:

- **Structured Scene Representations**: The images are parsed into detailed structural representations, capturing the property of the object, spatial relationships, and context of the scene. This approach reduces ambiguity and ensures fully interpretable reasoning processes.
- **Program Execution for Reasoning**: A deterministic program trace is associated with each query, enabling clear decision-making and accurate examination of model reasoning steps.
- **Efficient Learning**: NS-VQA systems can achieve near-perfect reasoning accuracy with significantly smaller training sets than traditional end-to-end neural models, demonstrating the efficiency of symbolic approaches.
- **Interpretability and Robustness**: The dataset's emphasis on symbolic reasoning provides strong interpretability, allowing models to generate clear, step-by-step explanations of their outputs, which is critical for high-stakes applications like healthcare and autonomous systems. [21]

The NS-VQA dataset highlights the advantages of structured, symbolic reasoning over purely neural approaches, reinforcing the importance of hybrid architectures in achieving human-like understanding.

D. Comparison and Emerging Trends

Together, these datasets provide a comprehensive foundation for evaluating NSAI systems, each addressing different aspects of the reasoning spectrum. While CLEVR focuses on controlled, synthetic reasoning, GQA extends this approach to real-world scenes, and NS-VQA emphasizes transparent, program-based reasoning. Ongoing research aims to further integrate these strengths, creating benchmarks that balance interpretability, scalability, and real-world applicability.

VIII. Future Directions for Neuro-Symbolic AI

By combining the structured reasoning of symbolic systems with the flexible learning of neural networks, neuro- symbolic artificial intelligence (NSAI) has become a potential paradigm for reaching human-like intelligence. But in order to fully utilise NSAI, a number of important issues must be resolved, and new lines of inquiry must be investigated. The main future directions that have been identified by current research are outlined in this part. These include unified representations, scalable learning, ethical considerations, improved explainability, practical applications, and new multidisciplinary techniques. [19] [13] [12] [22]

A. Unified Representations and Symbol Grounding

Creating unified representations that smoothly combine neural and symbolic components is one of the main problems in NSAI. The symbol grounding problem, which requires abstract symbols to be meaningfully connected to actual sensory inputs, is a common challenge for traditional NSAI systems. Zhang and Sheng (2024) emphasize the need for more coherent integration strategies that bridge the gap between low-level perceptual data and high-level symbolic reasoning. They suggest that future NSAI architectures should focus on hybrid frameworks that can dynamically adapt their internal representations based on context, improving both interpretability and scalability. [19]

Cunnington et al. (2024) further highlight the potential of foundation models, such as NeSyGPT, in addressing this challenge. These models leverage large-scale pre-trained representations to extract symbolic features from raw data, significantly reducing the need for extensive manual rule engineering. This approach has shown promise in improving the scalability and flexibility of NSAI systems, particularly in complex, multimodal tasks. Additionally, this integration can facilitate cross-domain transfer learning, allowing models trained in one context to be effectively adapted to new, unfamiliar environments. [13]

In order to bridge the gap between brain perception and symbolic logic, recent research also investigates the use of knowledge graphs and semantic embeddings, which allow models to create internal representations that are more precise and context-aware. The interpretability and resilience of NSAI systems could be greatly enhanced by this method, especially in practical situations where context is crucial. [19] [13]

B. Scalable Learning and Efficient Architectures

Scalability is still a major obstacle to the broad use of NSAIs. The high-dimensional data processing capabilities of neural networks and the combinatorial complexity of symbolic reasoning are difficult for many current architectures to balance. Wan et al. (2024) argue for the development of efficient, scalable NSAI architectures that can handle large, heterogeneous datasets while maintaining interpretability. They propose integrating hardware accelerators, optimized neural-symbolic pipelines, and energy-efficient learning algorithms to address this challenge. [12]

Bougzime et al. (2025) also emphasize the importance of reducing data requirements for NSAI systems. They argue that future models should focus on learning from smaller, domain-specific datasets while leveraging symbolic reasoning to generalize across diverse tasks. This approach not only reduces the computational footprint of NSAI systems but also enhances their adaptability to real-world environments. Additionally, emerging technologies like neuromorphic computing and quantum neural-symbolic systems hold promise for significantly increasing the speed and efficiency of NSAI systems. [22]

C. Enhanced Explainability and Interpretability

Explainability is a cornerstone of NSAI, critical for building trust in AI systems used in high-stakes applications like healthcare, finance, and autonomous systems. Zhang and Sheng (2024) propose that future NSAI research should focus on making intermediate representations more transparent, allowing both human users and downstream AI components to understand and validate decision processes. This includes developing new methods for tracing decision paths and generating human-readable explanations from complex neural-symbolic interactions. [19]

Cunnington et al. (2024) add that integrating foundation models with symbolic reasoning offers a promising path towards more interpretable NSAI systems. By leveraging pre-trained knowledge, these models can generate more structured, meaningful outputs, improving both transparency and reliability. Additionally, current research in neuro-symbolic causal reasoning has demonstrated promise in developing models that shed light on cause-and-effect links within complex systems in addition to explaining their decisions. [13]

D. Ethical and Societal Considerations

As NSAI systems become more capable, their ethical implications must be carefully considered. Bougzime et al. (2025) highlight the need for robust, ethically-aligned design principles that ensure AI systems operate safely and fairly in real-world contexts. This includes addressing potential biases, ensuring data privacy, and implementing fail-safe mechanisms to prevent unintended consequences. They also stress the importance of developing standardized benchmarks and evaluation frameworks to assess the ethical impact of NSAI technologies. [22]

Additionally, researchers are exploring the concept of "explainable accountability," where NSAI systems are required to provide detailed, auditable records of their decision-making processes to ensure compliance with ethical standards. This approach aims to create AI systems that are not only interpretable but also accountable for their actions, reducing the risk of misuse and unintended harm. [22]

E. Real-World Applications and Domain-Specific Advances

Finally, the practical deployment of NSAI systems will require significant advances in real-world applications. Wan et al. (2024) identify several promising areas, including collaborative robotics, human-AI interaction, and cognitive computing. They contend that in order to increase performance and dependability, NSAI systems need to be customised to meet the particular needs of every domain, utilising domain-specific expertise. [12]

As NSAI develops further, research must concentrate on resolving these issues and creating reliable, scalable, and morally good systems that can reason and comprehend at the level of humans. This will require a multidisciplinary approach, integrating insights from neuroscience, cognitive science, machine learning, and symbolic logic.

IX. Summary

Neuro-Symbolic AI (NSAI) represents a significant evolution in artificial intelligence, bridging the gap between the structured, rule-based reasoning of symbolic AI and the flexible, data-driven learning of neural networks. This hybrid approach seeks to combine the best of both worlds, offering systems that can understand, learn, and reason in ways that are both powerful and interpretable. Over the course of this literature review, we have explored the foundational concepts, historical evolution, key approaches, applications, challenges, and future directions of NSAI, highlighting the transformative potential of this emerging field.

The historical development of NSAI reflects the ongoing effort to overcome the limitations of purely neural and purely symbolic approaches. Early symbolic systems were transparent and interpretable, but because they couldn't generalise from raw data, they needed a lot of manual rule encoding. On the other hand, purely neural models frequently function as "black boxes" with little explainability, even when they have remarkable pattern recognition abilities. The emergence of NSAI addresses these gaps by integrating symbolic reasoning with neural learning, creating systems capable of both precise logic and adaptive learning.

Central to this integration are the foundational concepts of representation spaces, knowledge representation, and logic, which form the backbone of NSAI architectures. Representation spaces define how information is encoded and processed, enabling systems to balance structured logic with flexible, context-aware learning. In order for NSAI systems to produce human-readable justifications for intricate decisions, knowledge representation serves as the vital link between low-level neural activations and high-level symbolic reasoning. Logic, as a core component, supports structured reasoning, enabling these systems to capture and process abstract relationships in a human-understandable format.

We have also reviewed the major approaches and architectures within NSAI, including symbolic-neural integration, modular architectures, and end-to-end differentiable systems. The various approaches taken to integrate neural learning and symbolic logic are reflected in these designs, each of which is suited to particular learning, scalability, and reasoning difficulties. These methods, which range from Logic Tensor Networks to Neuro-Symbolic Concept Learners, demonstrate the continuous innovation in the field.

Applications of NSAI span a wide range of domains, including healthcare, geoscience, military systems, and patient monitoring, demonstrating its potential to transform industries that require both precise reasoning and adaptive learning. The use of specialized datasets and benchmarks, such as CLEVR, GQA, and NS-VQA, has played a critical role in advancing this field, providing rigorous testing environments that push the boundaries of AI capability.

Despite this progress, significant challenges remain. Scalability, knowledge acquisition, interpretability, and ethical considerations continue to limit the widespread adoption of NSAI systems. Addressing these challenges will require continued innovation in architecture design, learning algorithms, and data management, as well as a deeper understanding of human cognition and reasoning.

Looking ahead, the future of NSAI appears promising, with ongoing research focusing on unified representations, scalable learning, enhanced explainability, and ethically sound designs. As these systems become more capable, they are poised to play a critical role in the next generation of AI, bridging the gap between human-like understanding and machine intelligence.

In conclusion, neuro-Symbolic AI is leading the way in AI research and offers a path towards more adaptable, intelligent, and interpretable systems. This field has the potential to unleash new levels of AI power by combining the advantages of neural learning and symbolic reasoning, which could revolutionize industry and society in the years to come.

References

[1] B. C. Colelough and W. Regli, "Neuro-Symbolic AI in 2024: A Systematic Review," Apr. 2025, arXiv:2501.05435 [cs]. [Online]. Available: http://arxiv.org/abs/2501.05435

[2] A. Sheth, K. Roy, and M. Gaur, "Neurosymbolic AI – Why, What, and How," May 2023, arXiv:2305.00813 [cs]. [Online]. Available: http://arxiv.org/abs/2305.00813

[3] L. D. Raedt, S. Dumančić, R. Manhaeve, and G. Marra, "From Statistical Relational to Neuro-Symbolic Artificial Intelligence," Mar. 2020, arXiv:2003.08316 [cs]. [Online]. Available: http://arxiv.org/abs/2003.08316

[4] A. d. Garcez and L. C. Lamb, "Neurosymbolic AI: The 3rd Wave," Dec. 2020, arXiv:2012.05876 [cs]. [Online]. Available: http://arxiv.org/abs/2012.05876

[5] M. K. Sarker, L. Zhou, A. Eberhart, and P. Hitzler, "Neuro-Symbolic Artificial Intelligence: Current Trends," May 2021, arXiv:2105.05330 [cs]. [Online]. Available: http://arxiv.org/abs/2105.05330

[6] X. Zhang and V. S. Sheng, “Bridging the Gap: Representation Spaces in Neuro-Symbolic AI,” Nov. 2024, arXiv:2411.04393 [cs]. [Online]. Available: http://arxiv.org/abs/2411.04393

[7] I. Hatzilygeroudis and J. Prentzas, “Neuro-Symbolic Approaches for Knowledge Representation in Expert Systems,” *International Journal of Hybrid Intelligent Systems*, vol. 1, no. 3-4, pp. 111–126, Jan. 2005. [Online]. Available: https://journals.sagepub.com/doi/full/10.3233/HIS-2004-13-401

[8] V. Belle, “On the relevance of logic for AI, and the promise of neuro-symbolic learning.”

[9] L. Serafini and A. d. Garcez, “Logic Tensor Networks: Deep Learning and Logical Reasoning from Data and Knowledge,” Jul. 2016, arXiv:1606.04422 [cs]. [Online]. Available: http://arxiv.org/abs/1606.04422

[10] J. Johnson, B. Hariharan, L. v. d. Maaten, L. Fei-Fei, C. L. Zitnick, and R. Girshick, “CLEVR: A Diagnostic Dataset for Compositional Language and Elementary Visual Reasoning,” Dec. 2016, arXiv:1612.06890 [cs]. [Online]. Available: http://arxiv.org/abs/1612.06890

[11] D. A. Hudson and C. D. Manning, “GQA: A New Dataset for Real-World Visual Reasoning and Compositional Question Answering,” May 2019, arXiv:1902.09506 [cs]. [Online]. Available: http://arxiv.org/abs/1902.09506

[12] Z. Wan, C.-K. Liu, H. Yang, C. Li, H. You, Y. Fu, C. Wan, T. Krishna, Y. Lin, and A. Raychowdhury, “Towards Cognitive AI Systems: a Survey and Prospective on Neuro-Symbolic AI,” Jan. 2024, arXiv:2401.01040 [cs]. [Online]. Available: http://arxiv.org/abs/2401.01040

[13] D. Cunnington, M. Law, J. Lobo, and A. Russo, “The Role of Foundation Models in Neuro-Symbolic Learning and Reasoning,” Feb. 2024, arXiv:2402.01889 [cs]. [Online]. Available: http://arxiv.org/abs/2402.01889

[14] S. Paul, J. Yu, J. J. Dekker, A. Ignatiev, and P. J. Stuckey, “Formal Explanations for Neuro-Symbolic AI,” Oct. 2024, arXiv:2410.14219 [cs]. [Online]. Available: http://arxiv.org/abs/2410.14219

[15] Z. Wan, C.-K. Liu, H. Yang, R. Raj, C. Li, H. You, Y. Fu, C. Wan, S. Li, Y. Kim, A. Samajdar, Y. C. Lin, M. Ibrahim, J. M. Rabaey, T. Krishna, and A. Raychowdhury, “Towards Efficient Neuro-Symbolic AI: From Workload Characterization to Hardware Architecture,” Sep. 2024, arXiv:2409.13153 [cs]. [Online]. Available: http://arxiv.org/abs/2409.13153

[16] W. Chen, X. Ma, Z. Wang, W. Li, C. Fan, J. Zhang, X. Que, and C. Li, “Exploring neuro-symbolic AI applications in geoscience: implications and future directions for mineral prediction,” *Earth Science Informatics*, vol. 17, no. 3, pp. 1819–1835, Jun. 2024. [Online]. Available: https://link.springer.com/10.1007/s12145-024-01278-7

[17] D. H. Hagos and D. B. Rawat, “Neuro-Symbolic AI for Military Applications,” Aug. 2024, arXiv:2408.09224 [cs]. [Online]. Available: http://arxiv.org/abs/2408.09224

[18] O. Fenske, S. Bader, and T. Kirste, “Neuro-Symbolic Artificial Intelligence for Patient Monitoring,” Jun. 2024, arXiv:2407.11995 [cs]. [Online]. Available: http://arxiv.org/abs/2407.11995

[19] X. Zhang and V. S. Sheng, “Neuro-Symbolic AI: Explainability, Challenges, and Future Trends,” Nov. 2024, arXiv:2411.04383 [cs]. [Online]. Available: http://arxiv.org/abs/2411.04383

[20] D. Hossain and J. Y. Chen, “A Study on Neuro-Symbolic Artificial Intelligence: Healthcare Perspectives,” Mar. 2025, arXiv:2503.18213 [cs]. [Online]. Available: http://arxiv.org/abs/2503.18213

[21] K. Yi, J. Wu, C. Gan, A. Torralba, P. Kohli, and J. B. Tenenbaum, “Neural-Symbolic VQA: Disentangling Reasoning from Vision and Language Understanding,” Jan. 2019, arXiv:1810.02338 [cs]. [Online]. Available: http://arxiv.org/abs/1810.02338

[22] O. Bougzime, S. Jabbar, C. Cruz, and F. Demoly, “Unlocking the Potential of Generative AI through Neuro-Symbolic Architectures: Benefits and Limitations,” Feb. 2025, arXiv:2502.11269 [cs]. [Online]. Available: http://arxiv.org/abs/2502.11269